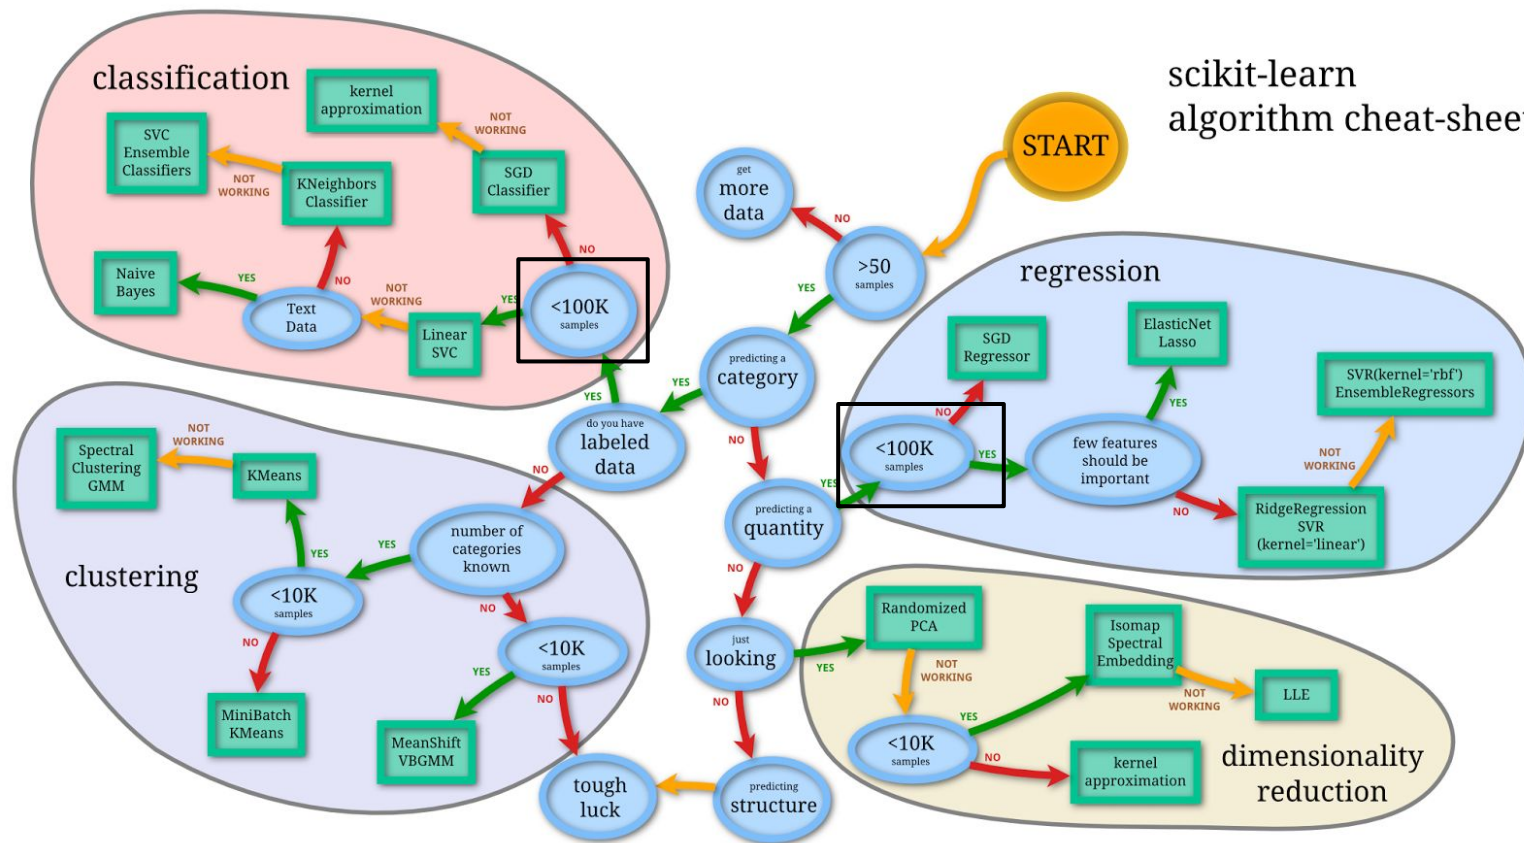

CNN

— Juan Luis García Mendoza, Mario Ezra
Aragón, Adrián Pastor López Monroy, Luis
Villaseñor Pineda, Manuel Montes y Gómez —

Overview

- Explaining How Convolutions Work
- Pooling
- CNN in NLP

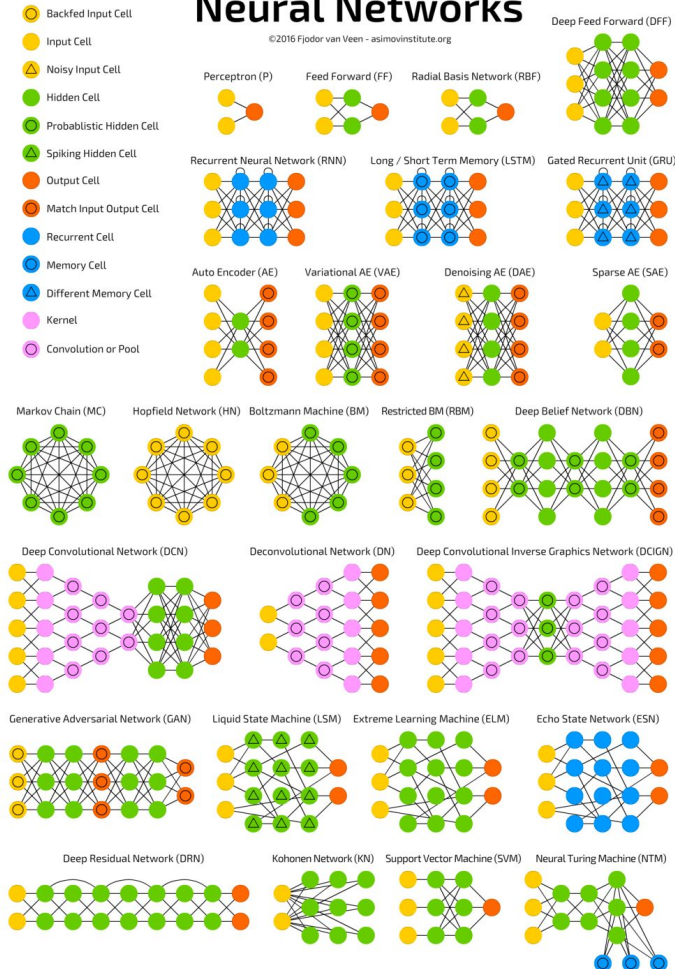

scikit-learn algorithm cheat-sheet



Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org


Deep Convolutional Network (DCN)




Terminology

- **filter size:** refers to the window size of the convolution operation
- **stride:** refers to the distance between two movements of the convolution window
- **padding:** refers to the way you handle boundaries of the input
- **pooling operation:** was introduced to CNNs mainly for reducing the size of the intermediate outputs.


Architecture



Architecture




Standard Convolution



$$h_{i,j} = \sum_{k=1}^m \sum_{l=1}^m w_{k,l} x_{i+k-1, j+l-1} \text{ where } 1 \leq i, j \leq n - m + 1$$

filter size (m) = 3, stride = 1, padding = None


Convolution with stride



$$h_{i,j} = \sum_{k=1}^m \sum_{l=1}^m w_{k,l} x_{(i-1) \times s_i + k, (j-1) \times s_j + l} \text{ where } 1 \leq i \leq \text{floor}[(n-m)/s_i] + 1 \text{ and } \text{floor}[(n-m)/s_j] + 1$$

filter size (m) = 2, stride = 2, padding = None


Convolution with padding



$$h_{i,j} = \sum_{k=1}^m \sum_{l=1}^m w_{k,l} x_{i+k-(m-1), j+l-(m-1)} \text{ where } 1 \leq i, j \leq n$$

filter size (m) = 3, stride = 1, padding = zero

Convolution (example)



filter size (m) = ?, stride = ?, padding = ?

1	1	0	0	1
0	0	0	0	1
0	1	1	1	1
0	0	1	0	1
1	0	1	1	1

$(1*1+0*1+1*0+0*0+1*0+1*0+1*0+1*1+0*1)=2$		

1	1	0	0	1
0	0	0	0	1
0	1	1	1	1
0	0	1	0	1
1	0	1	1	1

2	$(1*1+0*0+1*0+1*0+1*0+1*0+1*1+1*1+0*1)=3$	

1	1	0	0	1
0	0	0	0	1
0	1	1	1	1
0	0	1	0	1
1	0	1	1	1

2	3	$(1*0+0*0+1*1+0*0+1*0+1*0+1*1+1*1+0*1)=3$

1	1	0	0	1
0	0	0	0	1
0	1	1	1	1
0	0	1	0	1
1	0	1	1	1

2		3	3
$(1*0+0*0+1*0+0*0+1*1+0*1+1*0+1*0+1*1)=2$			

1	1	0	0	1
0	0	0	0	1
0	1	1	1	1
0	0	1	0	1
1	0	1	1	1


2	3		3
2	$(1*0+0*0+1*0+0*1+1*1+0*1+1*0+1*1+0*0)=2$		

Continuing to scan through input matrix and the final feature map will be


2	3	3
2	2	3
2	4	4

Activation functions


Tanh




ReLU




Sigmoid



Linear




Pooling operation



$$h_{i,j} = \max \left(\left\{ x_{i,j}, x_{i,j+1}, \dots, x_{i,j+m-1}, x_{i+1,j}, \dots, x_{i+1,j+m-1}, \dots, x_{i+m-1,j}, \dots, x_{i+m-1,j+m-1} \right\} \right) \text{ where } 1 \leq i, j \leq n - m + 1$$

filter size (m) = 3, stride = 1, padding = zero

Max pooling with stride




$$h_{i,j} = \max \left(\left\{ X_{(i-1) \times s_i + 1, (j-1) \times s_j + 1}, X_{(i-1) \times s_i + 1, (j-1) \times s_j + 2}, \dots, X_{(i-1) \times s_i + 1, (j-1) \times s_j + m}, X_{(i-1) \times s_i + 2, (j-1) \times s_j + 1}, \dots, X_{(i-1) \times s_i + 2, (j-1) \times s_j + m}, \dots, X_{(i-1) \times s_i + m, (j-1) \times s_j + 1}, \dots, X_{(i-1) \times s_i + m, (j-1) \times s_j + m} \right\} \right)$$

where $1 \leq i \leq \text{floor} \left[\frac{(n-m)}{s_i} \right] + 1$ and $1 \leq j \leq \text{floor} \left[\frac{(n-m)}{s_j} \right] + 1$

filter size (m) = 2, stride = 2, padding = zero


Average pooling



$$h_{i,j} = \frac{x_{i,j}, x_{i,j+1}, \dots, x_{i,j+m-1}, x_{i+1,j}, \dots, x_{i+1,j+m-1}, \dots, x_{i+m-1,j}, \dots, x_{i+m-1,j+m-1}}{m \times m} \quad \forall i \geq 1, j \leq n - m + 1$$

filter size (m) = 2, stride = 1, padding = zero

Example



Fully connected layers

- The initial **fully connected layer** found immediately after the **last convolution or pooling layer**.
- The weight matrix will be $w^{(m, h \times w \times d)}$.
- During inference (or prediction), we reshape the output of the last convolution/pooling layer to be of size $(h \times w \times d, 1)$
- The **output layer** can be a **softmax classification** layer for a classification problem or a **linear layer** for a regression problem.


CNN in NLP (input)

Let a sentence of **p** words and the length of the input sentences **n** words:

- We will pad the sentence with some special word (if **p < n**).
- We will represent each word in the sentence by a vector of size **k** (one-hot-encoded representation, word2vec, glove).
- A batch of sentences of size **b** can be represented by a **$b \times n \times k$** matrix.

CNN in NLP (convolution operation)

Image



6 x 6 x 3

$h \times w \times d$

Text (one sentence)

Bob
and
Mary
are
friends
PAD
PAD

1	0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0

7 x 13


$h \times w \rightarrow n \times k$

CNN in NLP (convolution operation)

By convolving the input \mathbf{x} ($\mathbf{n \times k}$) with a weight matrix \mathbf{W} ($\mathbf{m \times k}$), we will produce an output of \mathbf{h} ($\mathbf{1 \times n}$)

$$h_{i,l} = \sum_{j=1}^m \sum_{l=1}^k w_{j,l} x_{i+j-1,l} \quad h = W * x + b$$

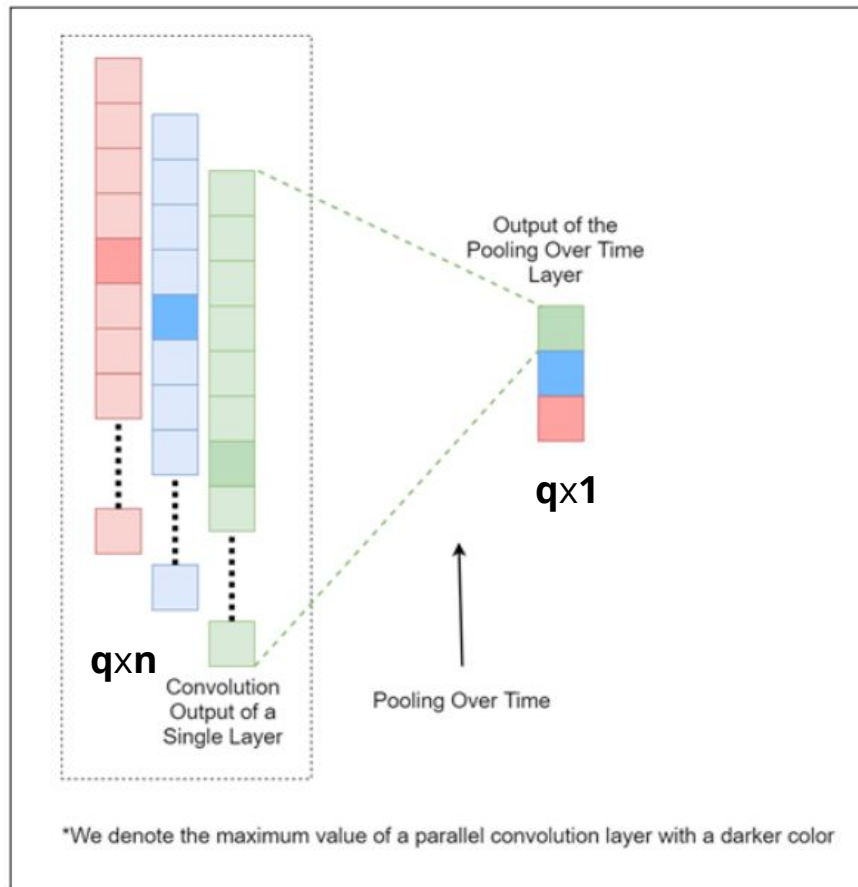
- We have parallel layers with different convolution filter sizes.
- Each convolution layer outputs a hidden vector of size $\mathbf{1 \times n}$.
- We will concatenate these outputs to form the input to the next layer of size $\mathbf{q \times n}$.
- \mathbf{q} is the number of parallel layers we will use.



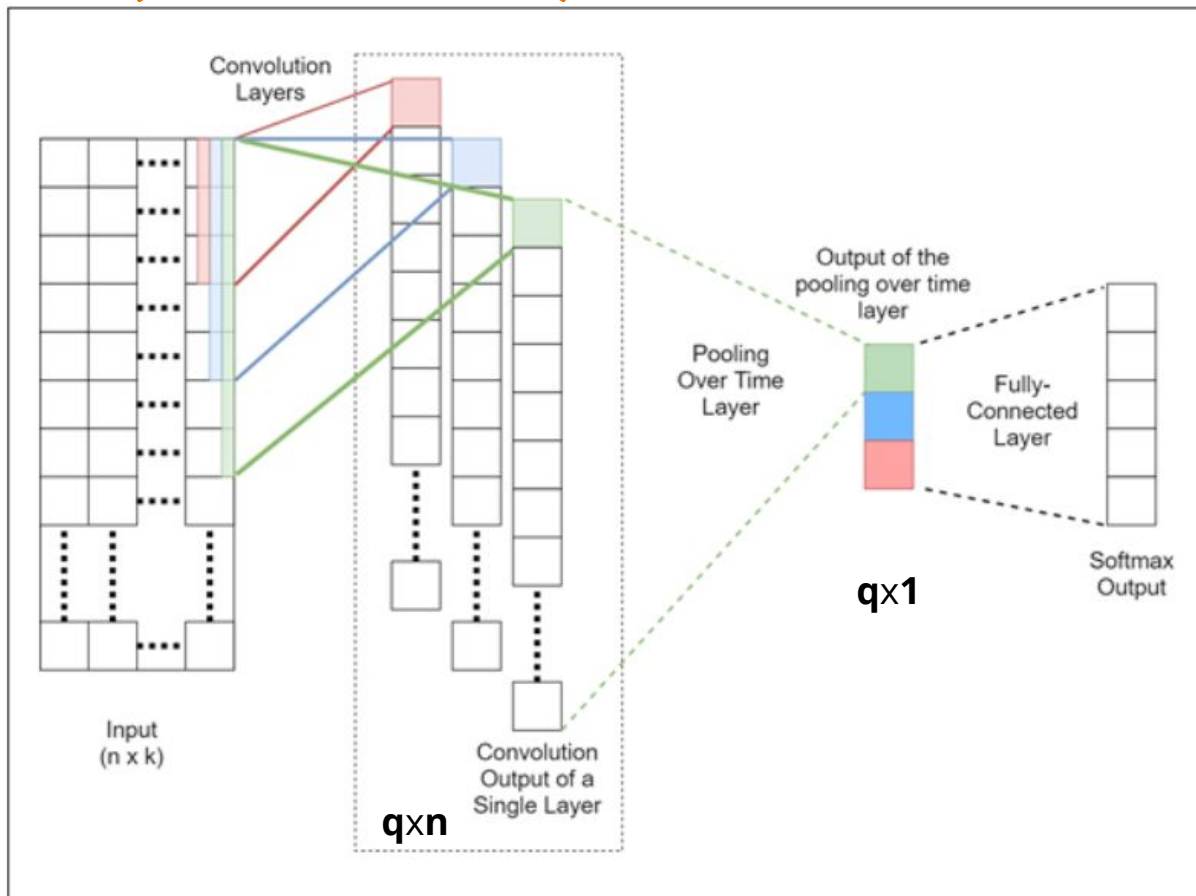
CNN in NLP (pooling operation)

- Let's assume the output of the last layer **h** is of size **qxn**.
- The pooling layer would produce an output **h'** of size **q×1** output.

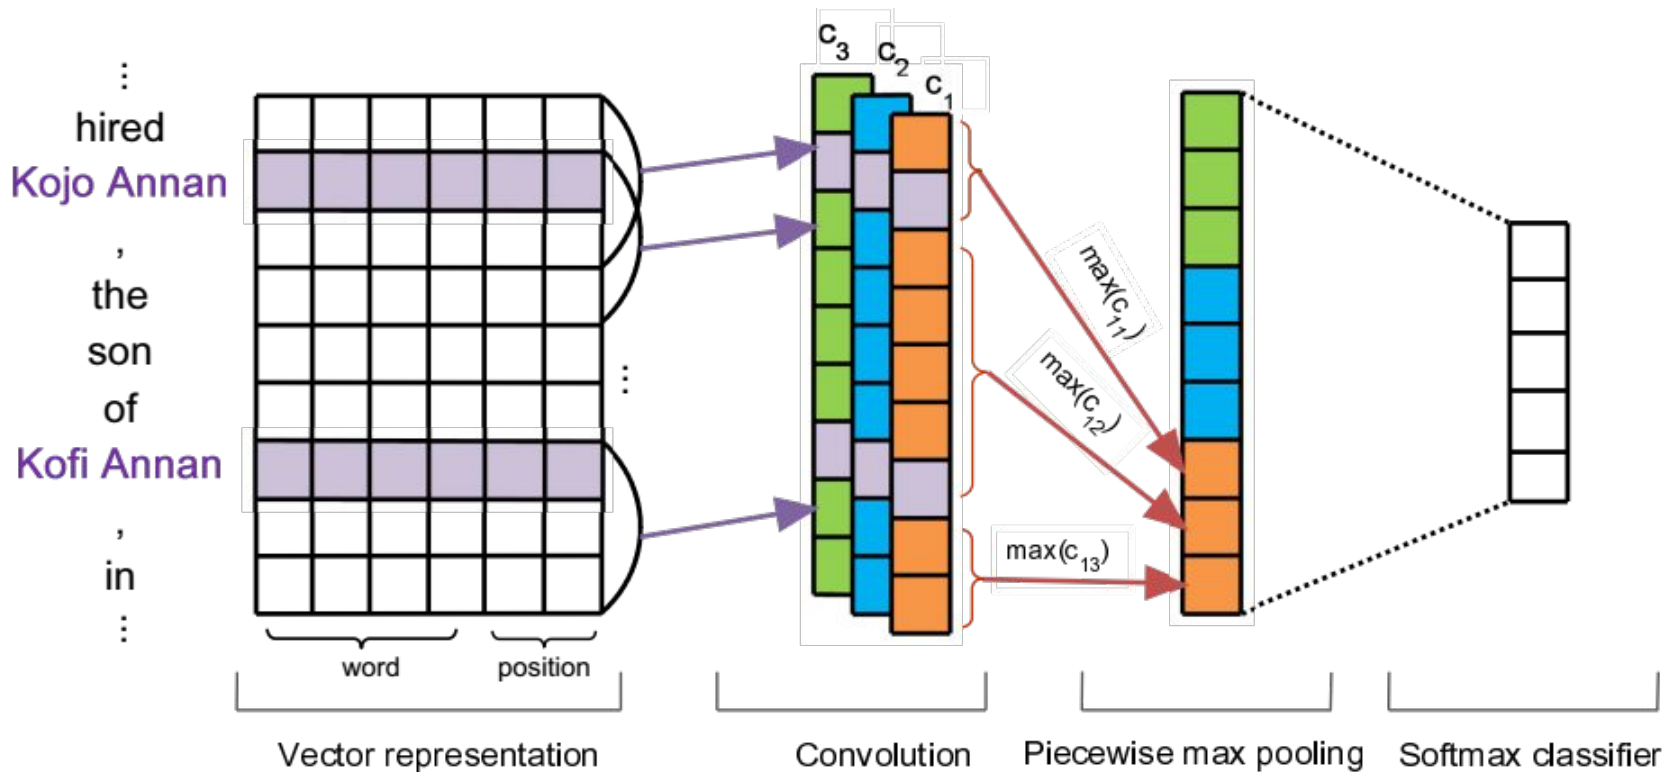
$$h'_{i,1} = \left\{ \max \left(h^{(i)} \right) \text{ where } 1 \leq i \leq q \right\}$$



CNN in NLP (architecture)



CNN in NLP (relation extraction)



References

- Ganegedara, T. (2018). Natural Language Processing with TensorFlow. Packt Publishing Ltd.
<https://github.com/PacktPublishing/Natural-Language-Processing-with-TensorFlow>
- Zeng, D., Liu, K., Chen, Y., & Zhao, J. (2015). Distant supervision for relation extraction via piecewise convolutional neural networks. In Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing (pp. 1753–1762).

Auxiliar Bibliography

- Stefan Kojouharov, Cheat Sheets for AI, Neural Networks, Machine Learning, Deep Learning & Big Data
- Renu Khandelwal, Convolutional Neural Network: Feature Map and Filter Visualization
- Renu Khandelwal, Convolutional Neural Network(CNN) Simplified
- Prabhu, Understanding of Convolutional Neural Network (CNN) — Deep Learning
- Activation Function

Questions?

Exercises

- Sentence Classification with CNN (Simple Architecture)
- Sentence Classification with CNN (Complex Architecture)

CNN

— Juan Luis García Mendoza, Mario Ezra
Aragón, Adrián Pastor López Monroy, Luis
Villaseñor Pineda, Manuel Montes y Gómez —
